

Pycaret is a low-code python library for ML

1 Imports

```
In [1]: import pandas as pd  
from pycaret.classification import *
```

2 Load your data (example)

```
In [3]: data = pd.read_csv("Thyroid_Diff.csv")
```

3 Setup the PyCaret experiment

```
In [5]: clf_setup = setup(  
    data=data,  
    target='Recurred',      # or your target column name  
    session_id=123,  
    fold=5  
)
```

	Description	Value
0	Session id	123
1	Target	Recurrent
2	Target type	Binary
3	Target mapping	No: 0, Yes: 1
4	Original data shape	(383, 17)
5	Transformed data shape	(383, 50)
6	Transformed train set shape	(268, 50)
7	Transformed test set shape	(115, 50)
8	Numeric features	1
9	Categorical features	15
10	Preprocess	True
11	Imputation type	simple
12	Numeric imputation	mean
13	Categorical imputation	mode
14	Maximum one-hot encoding	25
15	Encoding method	None
16	Fold Generator	StratifiedKFold
17	Fold Number	5
18	CPU Jobs	-1
19	Use GPU	False
20	Log Experiment	False
21	Experiment Name	clf-default-name

Description	Value
22	USI bfa0

4 Find the best model and store it in best_model

```
In [12]: from pycaret.classification import compare_models, tune_model, save_model

# 1) Find the best baseline model (already have this)
best_model = compare_models(sort='AUC')

# 2) Tune that best model, still optimizing for AUC
best_tuned_model = tune_model(best_model, optimize='AUC')

# 3) (Optional) save the tuned best model as a .pkl file
save_model(best_tuned_model, 'best_tuned_diabetes_model')
```

Model		Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
lightgbm	Light Gradient Boosting Machine	0.9626	0.9941	0.9626	0.9632	0.9622	0.9058	0.9074	0.3160
rf	Random Forest Classifier	0.9551	0.9907	0.9551	0.9552	0.9546	0.8868	0.8880	0.3460
et	Extra Trees Classifier	0.9551	0.9893	0.9551	0.9558	0.9546	0.8869	0.8889	0.3060
gbc	Gradient Boosting Classifier	0.9477	0.9889	0.9477	0.9481	0.9474	0.8698	0.8709	0.2740
lr	Logistic Regression	0.9476	0.9834	0.9476	0.9475	0.9465	0.8659	0.8681	0.1920
ridge	Ridge Classifier	0.9477	0.9751	0.9477	0.9484	0.9466	0.8660	0.8692	0.1680
svm	SVM - Linear Kernel	0.8769	0.9749	0.8769	0.9005	0.8743	0.6968	0.7252	0.1780
lda	Linear Discriminant Analysis	0.9477	0.9710	0.9477	0.9484	0.9466	0.8660	0.8692	0.1780
ada	Ada Boost Classifier	0.9440	0.9674	0.9440	0.9451	0.9436	0.8603	0.8626	0.2600
nb	Naive Bayes	0.9253	0.9402	0.9253	0.9279	0.9239	0.8105	0.8162	0.1860
qda	Quadratic Discriminant Analysis	0.8807	0.9307	0.8807	0.8862	0.8692	0.6645	0.6914	0.1940
dt	Decision Tree Classifier	0.9402	0.9260	0.9402	0.9425	0.9399	0.8511	0.8543	0.1800
knn	K Neighbors Classifier	0.8693	0.9095	0.8693	0.8769	0.8618	0.6502	0.6728	0.2600
dummy	Dummy Classifier	0.7164	0.5000	0.7164	0.5133	0.5981	0.0000	0.0000	0.1820

	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC
Fold							
0	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000
1	0.9444	0.9795	0.9444	0.9440	0.9438	0.8586	0.8596
2	0.9630	0.9967	0.9630	0.9630	0.9630	0.9112	0.9112
3	0.9434	0.9684	0.9434	0.9430	0.9428	0.8577	0.8586
4	0.9811	0.9956	0.9811	0.9816	0.9809	0.9526	0.9536
Mean	0.9664	0.9880	0.9664	0.9663	0.9661	0.9160	0.9166
Std	0.0218	0.0121	0.0218	0.0220	0.0220	0.0550	0.0547

Fitting 5 folds for each of 10 candidates, totalling 50 fits

Original model was better than the tuned model, hence it will be returned. NOTE: The display metrics are for the tuned model (not the original one).

Transformation Pipeline and Model Successfully Saved

```
Out[12]: (Pipeline(memory=Memory(location=None),
      steps=[('label_encoding',
               TransformerWrapperWithInverse(exclude=None, include=None,
                                              transformer=LabelEncoder())),
      ('numerical_imputer',
       TransformerWrapper(exclude=None, include=['Age'],
                          transformer=SimpleImputer(add_indicator=False,
                                                     copy=True,
                                                     fill_value=None,
                                                     keep_empty_features=False,
                                                     missing_values=np.nan,
                                                     strategy='mean'))...
      LGBMClassifier(boosting_type='gbdt', class_weight=None,
                     colsample_bytree=1.0, importance_type='split',
                     learning_rate=0.1, max_depth=-1,
                     min_child_samples=20, min_child_weight=0.001,
                     min_split_gain=0.0, n_estimators=100, n_jobs=-1,
                     num_leaves=31, objective=None, random_state=123,
                     reg_alpha=0.0, reg_lambda=0.0, subsample=1.0,
                     subsample_for_bin=200000, subsample_freq=0))],
      verbose=False),
      'best_tuned_diabetes_model.pkl')
```

```
In [13]: models()
```

Out[13]:

ID	Name	Reference	Turbo
lr	Logistic Regression	sklearn.linear_model._logistic.LogisticRegression	True
knn	K Neighbors Classifier	sklearn.neighbors._classification.KNeighborsCl...	True
nb	Naive Bayes	sklearn.naive_bayes.GaussianNB	True
dt	Decision Tree Classifier	sklearn.tree._classes.DecisionTreeClassifier	True
svm	SVM - Linear Kernel	sklearn.linear_model._stochastic_gradient.SGDC...	True
rbfsvm	SVM - Radial Kernel	sklearn.svm._classes.SVC	False
gpc	Gaussian Process Classifier	sklearn.gaussian_process._gpc.GaussianProcessC...	False
mlp	MLP Classifier	sklearn.neural_network._multilayer_perceptron....	False
ridge	Ridge Classifier	sklearn.linear_model._ridge.RidgeClassifier	True
rf	Random Forest Classifier	sklearn.ensemble._forest.RandomForestClassifier	True
qda	Quadratic Discriminant Analysis	sklearn.discriminant_analysis.QuadraticDiscrim...	True
ada	Ada Boost Classifier	sklearn.ensemble._weight_boosting.AdaBoostClas...	True
 gbc	Gradient Boosting Classifier	sklearn.ensemble._gb.GradientBoostingClassifier	True
 lda	Linear Discriminant Analysis	sklearn.discriminant_analysis.LinearDiscrimina...	True
 et	Extra Trees Classifier	sklearn.ensemble._forest.ExtraTreesClassifier	True
lightgbm	Light Gradient Boosting Machine	lightgbm.sklearn.LGBMClassifier	True
dummy	Dummy Classifier	sklearn.dummy.DummyClassifier	True


In [14]:

```
print(model_ridge_classifier)
```

```
RidgeClassifier(alpha=1.0, class_weight=None, copy_X=True, fit_intercept=True,
               max_iter=None, positive=False, random_state=123, solver='auto',
               tol=0.0001)
```

```
In [15]: # Confusion matrix
```


```
plot_model(tune_ridge_classifier, plot='confusion_matrix')
```



```
In [16]: # Classification error
```


```
plot_model(tune_ridge_classifier, plot='error')
```

Class Prediction Error for RidgeClassifier




```
In [17]: # Classification report (precision, recall, F1, etc.)  
plot_model(tune_ridge_classifier, plot='class_report')
```

RidgeClassifier Classification Report




```
In [18]: # Feature importance  
plot_model(tune_ridge_classifier, plot='feature')
```



```
In [19]: # Learning curve  
plot_model(tune_ridge_classifier, plot='learning')
```

Learning Curve for RidgeClassifier



```
In [20]: evaluate_model(tune_ridge_classifier)
```

```
interactive(children=(ToggleButtons(description='Plot Type:', icons=''), options=([('Pipeline Plot', 'pipeline'), ('Cross Validation Plot', 'cv')], value='cv')))
```

```
In [ ]:
```